


```
In [1]: ▶ import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm

# Load and Reformat Data
data = pd.read_csv('C:/Users/user/Desktop/My learning/ClinSoft/audi.csv')

# Data Exploration
print(data.columns)
print(data.info())
print(data.describe())
```

```

Index(['model', 'year', 'price', 'transmission', 'mileage', 'fuelType',
      'tax',
      'mpg', 'engineSize'],
      dtype='object')
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10668 entries, 0 to 10667
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   model           10668 non-null  object
1   year            10668 non-null  int64
2   price           10668 non-null  int64
3   transmission    10668 non-null  object
4   mileage         10668 non-null  int64
5   fuelType       10668 non-null  object
6   tax             10668 non-null  int64
7   mpg            10668 non-null  float64
8   engineSize     10668 non-null  float64
dtypes: float64(2), int64(4), object(3)
memory usage: 750.2+ KB
None

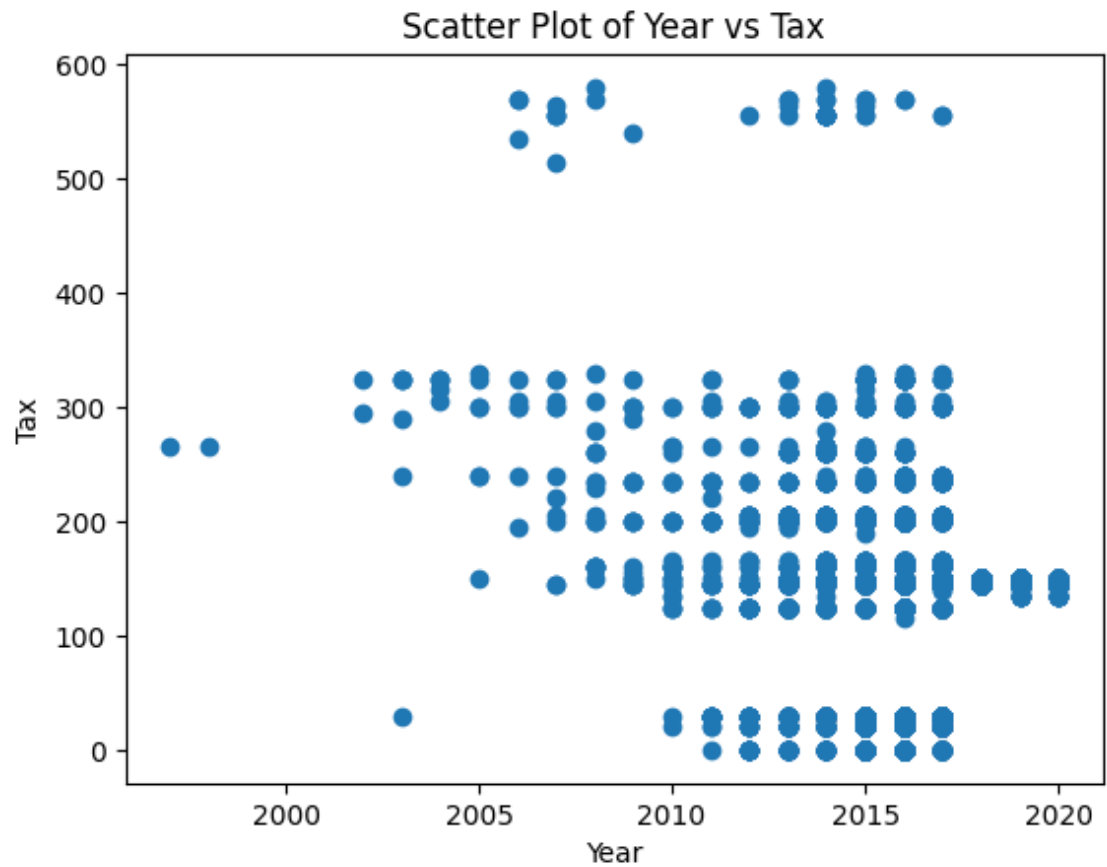
```

	year	price	mileage	tax	mpg
count	10668.000000	10668.000000	10668.000000	10668.000000	10668.000000
mean	2017.100675	22896.685039	24827.244001	126.011436	50.770022
std	2.167494	11714.841888	23505.257205	67.170294	12.949782
min	1997.000000	1490.000000	1.000000	0.000000	18.900000
25%	2016.000000	15130.750000	5968.750000	125.000000	40.900000
50%	2017.000000	20200.000000	19000.000000	145.000000	49.600000
75%	2019.000000	27990.000000	36464.500000	145.000000	58.900000
max	2020.000000	145000.000000	323000.000000	580.000000	188.300000

	engineSize
count	10668.000000
mean	1.930709
std	0.602957
min	0.000000
25%	1.500000
50%	2.000000
75%	2.000000
max	6.300000


In [2]:

```
# Scatter Plot
plt.scatter(data['year'], data['tax'])
plt.xlabel('Year')
plt.ylabel('Tax')
plt.title('Scatter Plot of Year vs Tax')
plt.show()
```



In [3]:


```
# Scatter Plot with Density  
sns.scatterplot(x='price', y='mileage', hue='fuelType', data=data, alpha=0.5)  
plt.show()
```



In [4]:

Density Plot


```
sns.kdeplot(data=data, x='price', hue='fuelType', alpha=0.3, fill=True)  
plt.show()
```



In [5]:

Histogram

```
sns.histplot(data=data, x='price', hue='fuelType', alpha=0.3, multiple='stack')  
plt.show()
```




In [6]:

```

# Filtering and Plotting
filtered_data = data[data['mileage'] < 180000]
sns.scatterplot(data=filtered_data, x='mileage', y='price', hue='transmiss
plt.show()

# Log Transform and Plot
filtered_data['log_price'] = np.log(filtered_data['price'])
sns.scatterplot(data=filtered_data, x='mileage', y='log_price', hue='trans
plt.show()

```




C:\Users\user\AppData\Local\Temp\ipykernel_4512\2771945094.py:7: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
filtered_data['log_price'] = np.log(filtered_data['price'])
```

In [7]:



```
# Linear Regression
model = sm.OLS.from_formula('log_price ~ mileage + tax', data=filtered_data)
result = model.fit()
print(result.summary())
```

OLS Regression Results

```

=====
=====
Dep. Variable:          log_price    R-squared:
0.555
Model:                  OLS          Adj. R-squared:
0.555
Method:                 Least Squares    F-statistic:
6645.
Date:                   Sat, 05 Aug 2023    Prob (F-statistic):
0.00
Time:                   22:53:06          Log-Likelihood:           -2
776.2
No. Observations:      10667          AIC:
5558.
Df Residuals:          10664          BIC:
5580.
Df Model:               2
Covariance Type:       nonrobust
=====
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	9.9745	0.008	1288.946	0.000	9.959	
mileage	-1.271e-05	1.32e-07	-96.116	0.000	-1.3e-05	-1.24e-05
tax	0.0021	4.59e-05	46.473	0.000	0.002	

```

=====
=====
Omnibus:                834.078    Durbin-Watson:
1.706
Prob(Omnibus):          0.000    Jarque-Bera (JB):           240
0.716
Skew:                   0.422    Prob(JB):
0.00
Kurtosis:               5.165    Cond. No.                8.6
7e+04
=====
=====


```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 8.67e+04. This might indicate that there are strong multicollinearity or other numerical problems.


In [9]:

```
# Residual Analysis
fig, ax = plt.subplots(figsize=(8, 6))
sns.residplot(x=result.fittedvalues, y=result.resid, lowess=True, line_kws=
plt.xlabel('Fitted Values')
plt.ylabel('Residuals')
plt.title('Residual Plot')
plt.show()
```



In [10]:


```
# Histogram of Residuals  
plt.hist(result.resid, bins=30, edgecolor='k', alpha=0.7)  
plt.xlabel('Residuals')  
plt.ylabel('Frequency')  
plt.title('Histogram of Residuals')  
plt.show()
```



In [11]:

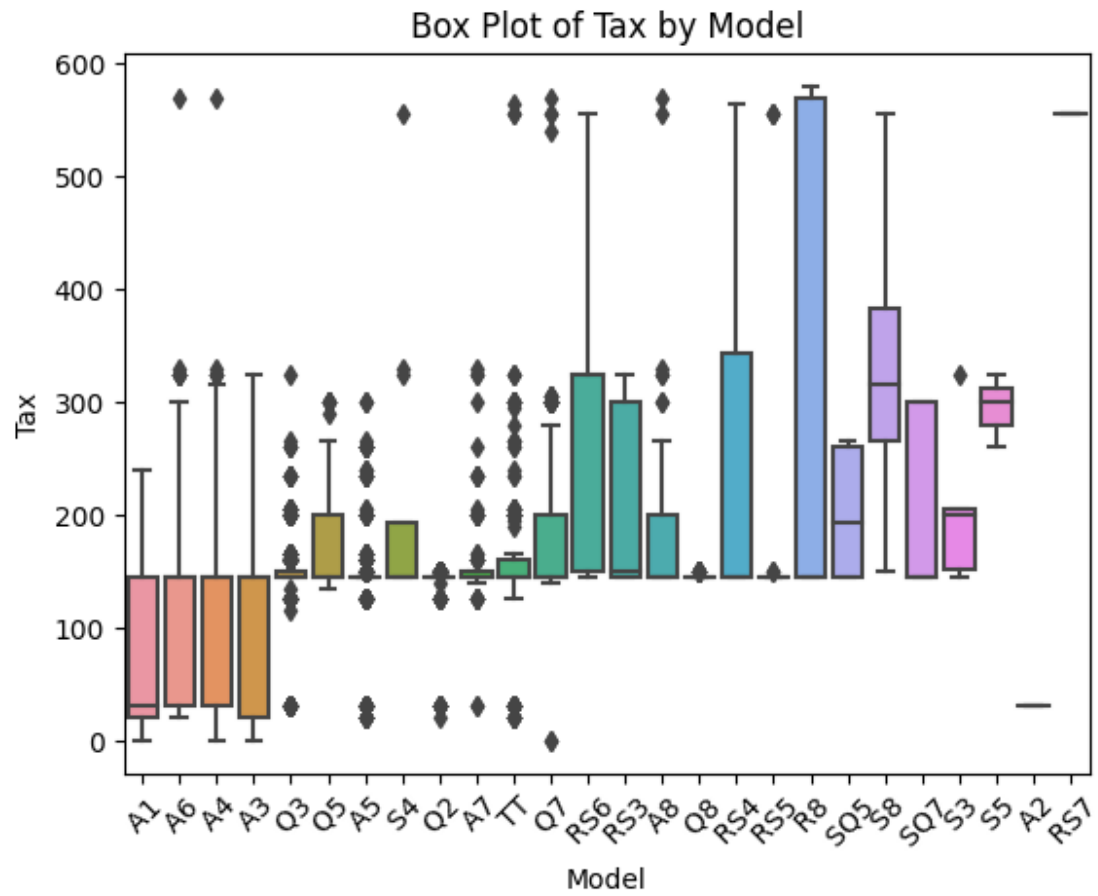
Filtering Outliers

```
filtered_data = data[(data['mileage'] < 180000) & (data['price'] > 20000)]  
sns.scatterplot(data=filtered_data, x='mileage', y=1/filtered_data['price'])  
plt.show()
```




In [12]:

```
# Box Plot
sns.boxplot(data=data, x='model', y='tax')
plt.xticks(rotation=45)
plt.xlabel('Model')
plt.ylabel('Tax')
plt.title('Box Plot of Tax by Model')
plt.show()
```



In [13]:

```
# Day 2
plt.scatter(data['mpg'], data['price'])
plt.xlabel('MPG')
plt.ylabel('Price')
plt.title('Scatter Plot of MPG vs Price')
plt.show()
```




```
In [19]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm


filtered_data['log_price'] = np.log(filtered_data['price'])
sns.scatterplot(data=filtered_data, x='mileage', y='log_price', hue='transmission')
plt.show()
```

C:\Users\user\AppData\Local\Temp\ipykernel_4512\969365871.py:9: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

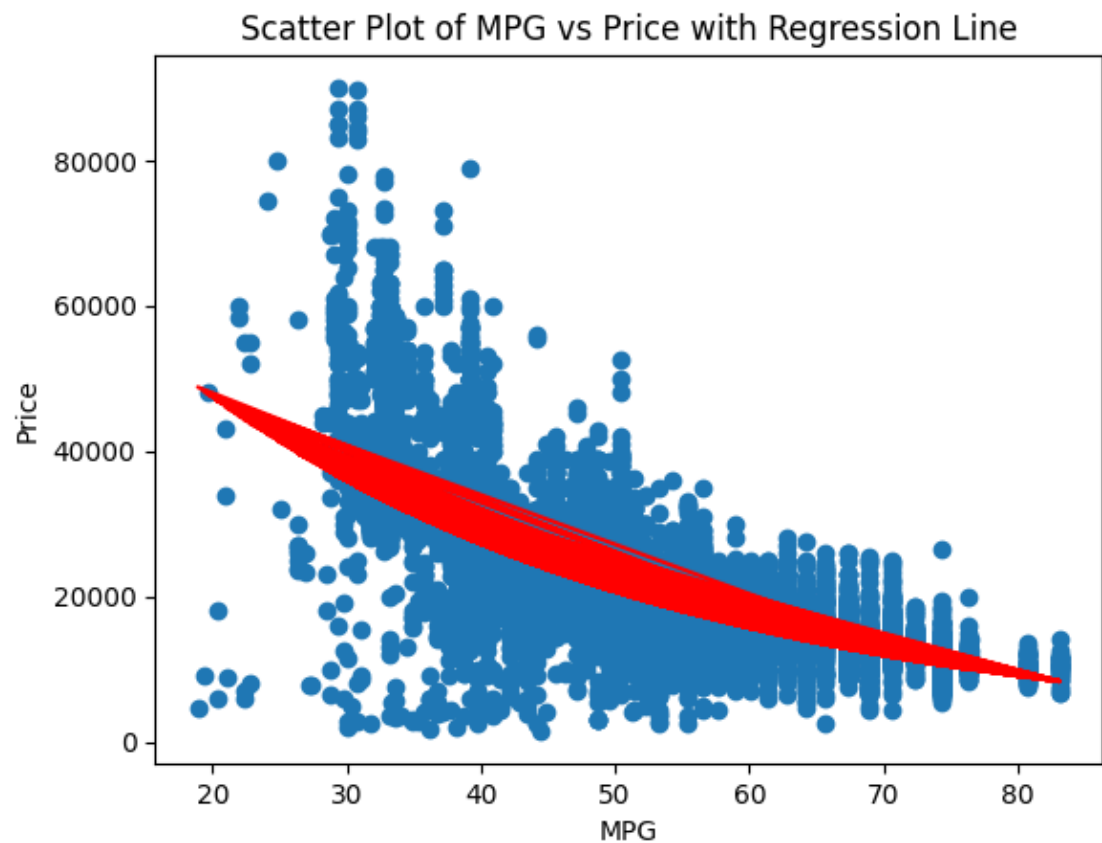
```
filtered_data['log_price'] = np.log(filtered_data['price'])
```



```
In [24]: > filtered_data2 = data[(data['mpg'] < 100) & (data['price'] < 90000)]

# Linear Regression with Log Transformed Price
model2 = sm.OLS.from_formula('np.log(price) ~ mpg', data=filtered_data2)
result2 = model2.fit()

plt.scatter(filtered_data2['mpg'], filtered_data2['price'])
plt.plot(filtered_data2['mpg'], np.exp(result2.fittedvalues), color='red')
plt.xlabel('MPG')
plt.ylabel('Price')
plt.title('Scatter Plot of MPG vs Price with Regression Line')
plt.show()
```




```
In [31]: filtered_data2 = data[(data['mpg'] < 100) & (data['price'] < 90000)]

# GLM with Log Transformed MPG
model3 = sm.GLM.from_formula('np.log(mpg) ~ price', data=filtered_data2, family=sm.families.Poisson)
result3 = model3.fit()

# Calculate Poisson Deviance Residuals
residuals = result3.resid_deviance

# Plot Residuals
fig, ax = plt.subplots(figsize=(8, 6))
sns.residplot(x=np.exp(result3.fittedvalues), y=residuals, lowess=True, linecolor='red')
plt.xlabel('Fitted Values')
plt.ylabel('Deviance Residuals')
plt.title('Residual Plot for Poisson GLM')
plt.show()
```




```
In [33]: # Removing Outliers
filtered_data4 = filtered_data2[filtered_data2['price'] < 100000]

# Linear Regression with Log Transformed MPG
model4 = sm.OLS.from_formula('np.log(mpg) ~ price', data=filtered_data4)
result4 = model4.fit()

# Calculate Linear Regression Residuals
residuals_lr = result4.resid

# Plot Residuals
fig, ax = plt.subplots(figsize=(8, 6))
sns.residplot(x=result4.fittedvalues, y=residuals_lr, lowess=True, line_kw=
plt.xlabel('Fitted Values')
plt.ylabel('Residuals')
plt.title('Residual Plot for Linear Regression')
plt.show()
```



```
In [ ]: 
```