


```
In [1]: ▶ import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm

# Read the data
data = pd.read_csv("C:/Users/user/Desktop/My learning/ClinSoft/audi.csv")

# Basic data exploration
print(data.columns)
print(data.info())
print(data.describe())
```

```

Index(['model', 'year', 'price', 'transmission', 'mileage', 'fuelType',
      'tax',
      'mpg', 'engineSize'],
      dtype='object')
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10668 entries, 0 to 10667
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   model                 10668 non-null  object
1   year                 10668 non-null  int64
2   price                10668 non-null  int64
3   transmission         10668 non-null  object
4   mileage              10668 non-null  int64
5   fuelType             10668 non-null  object
6   tax                  10668 non-null  int64
7   mpg                  10668 non-null  float64
8   engineSize           10668 non-null  float64
dtypes: float64(2), int64(4), object(3)
memory usage: 750.2+ KB
None

```

	year	price	mileage	tax	mpg
count	10668.000000	10668.000000	10668.000000	10668.000000	10668.000000
mean	2017.100675	22896.685039	24827.244001	126.011436	50.770022
std	2.167494	11714.841888	23505.257205	67.170294	12.949782
min	1997.000000	1490.000000	1.000000	0.000000	18.900000
25%	2016.000000	15130.750000	5968.750000	125.000000	40.900000
50%	2017.000000	20200.000000	19000.000000	145.000000	49.600000
75%	2019.000000	27990.000000	36464.500000	145.000000	58.900000
max	2020.000000	145000.000000	323000.000000	580.000000	188.300000

	engineSize
count	10668.000000
mean	1.930709
std	0.602957
min	0.000000
25%	1.500000
50%	2.000000
75%	2.000000
max	6.300000


In [2]:



```
# Scatter plot
plt.scatter(data['year'], data['tax'])
plt.xlabel('Year')
plt.ylabel('Tax')
plt.show()


# Scatter plot with color
sns.scatterplot(data=data, x='price', y='mileage', hue='fuelType', alpha=0.3)
plt.show()

# Density plot
sns.kdeplot(data=data, x='price', hue='fuelType', alpha=0.3)
plt.show()

# Histogram
sns.histplot(data=data, x='price', hue='fuelType', alpha=0.3)
plt.show()
```








In [3]:

```
# Filter and scatter plot  
first_data = data[data['mileage'] < 180000]  
sns.scatterplot(data=first_data, x='mileage', y='price', hue='transmission')  
plt.show()
```



In [4]: ▶

```
# Linear regression and residual analysis
model1 = sm.OLS(first_data['price'], sm.add_constant(first_data['mileage']))
print(model1.summary())
sm.graphics.plot_regress_exog(model1, 'mileage', fig=plt.figure(figsize=(10, 5)))
plt.show()
```


OLS Regression Results

```

=====
=====
Dep. Variable:          price    R-squared:
0.289
Model:                  OLS      Adj. R-squared:
0.289
Method:                 Least Squares    F-statistic:
4331.
Date:                   Sun, 06 Aug 2023    Prob (F-statistic):
0.00
Time:                   14:57:33    Log-Likelihood:          -1.132
5e+05
No. Observations:      10667    AIC:          2.26
5e+05
Df Residuals:          10665    BIC:          2.26
5e+05
Df Model:               1
Covariance Type:       nonrobust
=====
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	2.959e+04	139.601	211.967	0.000	2.93e+04	2.99e+04
mileage	-0.2699	0.004	-65.814	0.000	-0.278	-0.262

```

=====
=====
Omnibus:                7262.777    Durbin-Watson:
1.762
Prob(Omnibus):          0.000    Jarque-Bera (JB):          16597
7.135
Skew:                   2.958    Prob(JB):
0.00
Kurtosis:               21.397    Cond. No.          4.9
7e+04
=====
=====

```

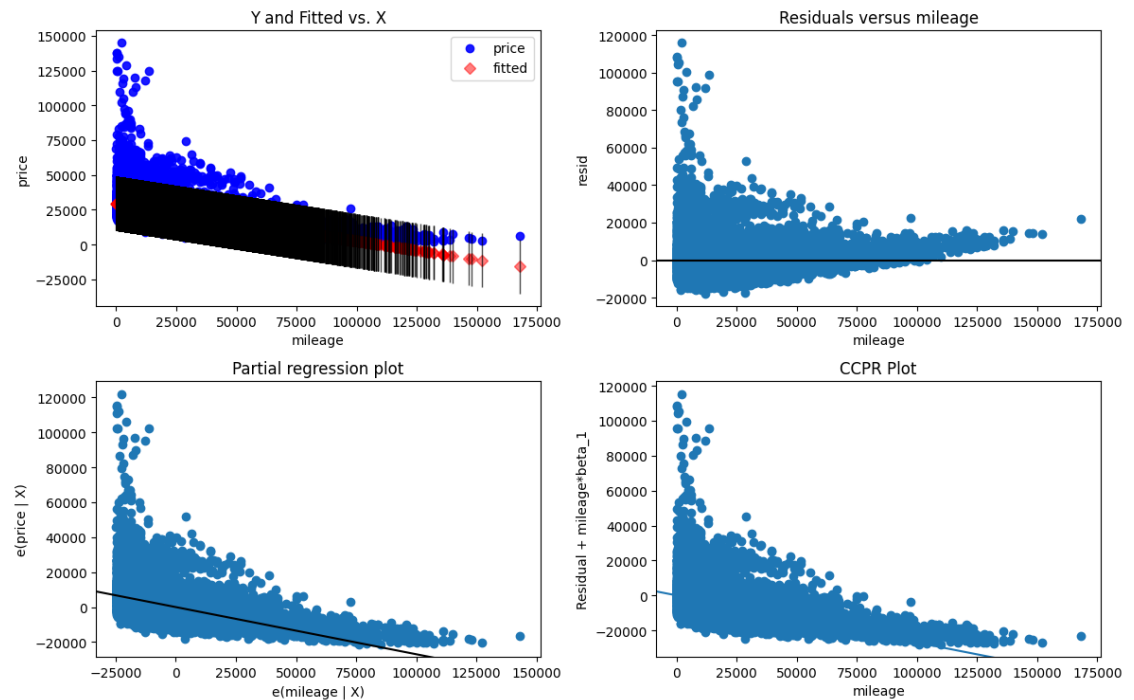
Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.97e+04. This might indicate that there are strong multicollinearity or other numerical problems.

eval_env: 1


Regression Plots for mileage



In [5]:



```
# Filter and scatter plot with another variable
second_data = data[(data['mileage'] < 180000) & (data['mpg'] < 100)]
sns.scatterplot(data=second_data, x='mpg', y='mileage')
plt.show()
```



In [6]:


```
# Linear regression with scatter plot and regression line
model2 = sm.OLS(second_data['mileage'], sm.add_constant(second_data['mpg']))
print(model2.summary())
sns.regplot(data=second_data, x='mpg', y='mileage')
plt.show()
```

OLS Regression Results

```
=====
=====
Dep. Variable:          mileage    R-squared:
0.189
Model:                  OLS        Adj. R-squared:
0.189
Method:                 Least Squares    F-statistic:
2474.
Date:                   Sun, 06 Aug 2023    Prob (F-statistic):
0.00
Time:                   14:58:07    Log-Likelihood:          -1.209
2e+05
No. Observations:       10634    AIC:                2.41
8e+05
Df Residuals:           10632    BIC:                2.41
8e+05
Df Model:                1
Covariance Type:        nonrobust
=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
const      -1.898e+04    903.199    -21.011    0.000    -2.07e+04    -1.7
2e+04
mpg          867.2376     17.435     49.740    0.000     833.061     90
1.414
=====
=====
Omnibus:                3768.420    Durbin-Watson:
1.289
Prob(Omnibus):           0.000    Jarque-Bera (JB):          1401
8.667
Skew:                    1.766    Prob(JB):
0.00
Kurtosis:                7.378    Cond. No.
230.
=====
=====
```


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.



In [7]:

```
# Boxplot
sns.boxplot(data=data, x='model', y='tax')
plt.xticks(rotation=90)
plt.show()
```



In [11]:

```
# Transformation and filtering
data['good_mpg'] = np.where((data['mpg'] >= 50) & (data['mpg'] <= 60), 'Good', 'Bad')
print(data.head())

# Linear regression on transformed data
model3 = sm.OLS(np.log(data['price']), sm.add_constant(data['mpg'])).fit()
print(model3.summary())
sm.graphics.plot_regress_exog(model3, 'mpg', fig=plt.figure(figsize=(12, 8)))
plt.show()
```

```

    model year price transmission mileage fuelType tax mpg engineSize
0   A1  2017  12500      Manual   15735   Petrol  150  55.4         1.
4   A6  2016  16500    Automatic   36203   Diesel   20  64.2         2.
0   A1  2016  11000      Manual   29946   Petrol   30  55.4         1.
4   A4  2017  16800    Automatic   25952   Diesel  145  67.3         2.
0   A3  2019  17300      Manual    1998   Petrol  145  49.6         1.
0

```

```

    good_mpg
0 Good MPG
1
2 Good MPG
3
4

```

OLS Regression Results

```

=====
=====
Dep. Variable:          price   R-squared:
0.379
Model:                OLS   Adj. R-squared:
0.379
Method:             Least Squares   F-statistic:
6519.
Date:              Sun, 06 Aug 2023   Prob (F-statistic):
0.00
Time:              15:00:53   Log-Likelihood:          -4
558.2
No. Observations:      10668   AIC:
9120.
Df Residuals:          10666   BIC:
9135.
Df Model:              1
Covariance Type:      nonrobust
=====
=====

```

```

=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
const          11.0651      0.015    761.339      0.000      11.037      1
1.094
mpg           -0.0224      0.000   -80.742      0.000     -0.023     -
0.022
=====
=====

```

```

=====
=====
Omnibus:              2062.976   Durbin-Watson:
1.395
Prob(Omnibus):         0.000   Jarque-Bera (JB):          5689
7.534
Skew:                 -0.200   Prob(JB):
0.00
Kurtosis:             14.307   Cond. No.

```

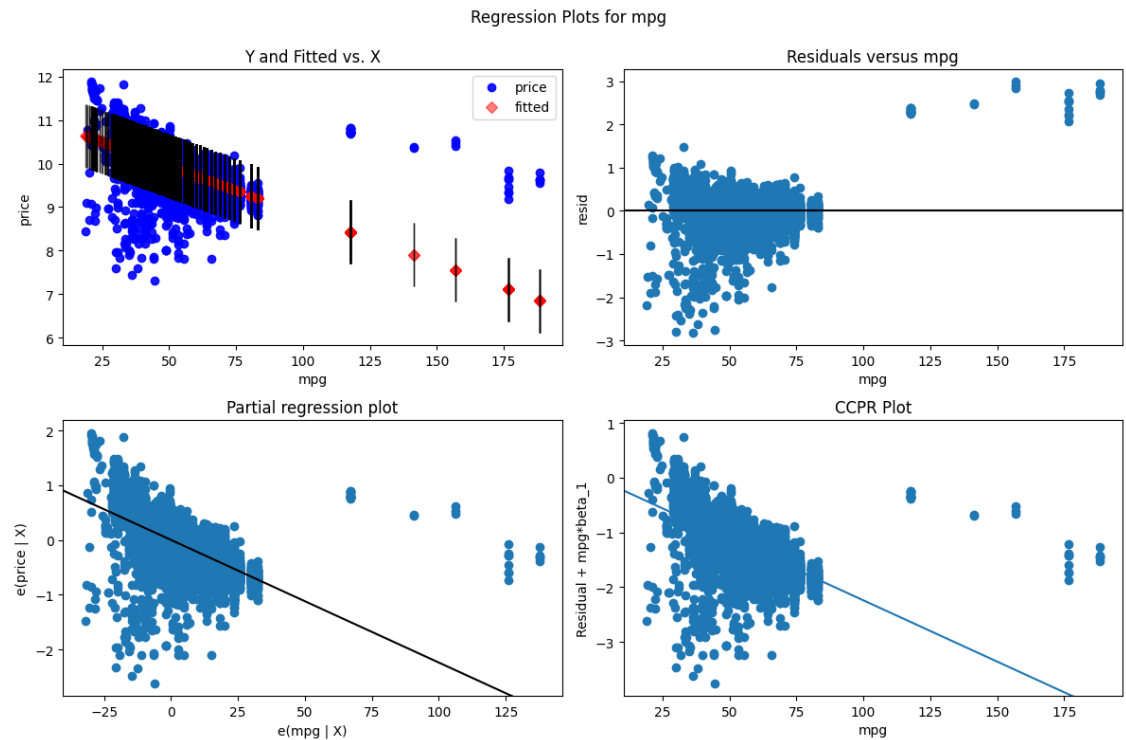
212.

```
=====
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


eval_env: 1



In [15]:

```
# Filtering and regression plot
filtered_data = data[(data['price'] < 100000) & (np.log(data['mpg']) < 4.5)
sns.lmplot(data=filtered_data, x='mpg', y='price', line_kws={'color': 'red'})
plt.show()
```

C:\Users\user\Anaconda3\lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight
self._figure.tight_layout(*args, **kwargs)



In [20]:

```
# Linear regression and residual analysis
model4 = sm.OLS(filtered_data['mpg'], sm.add_constant(filtered_data['price']))
print(model4.summary())
sm.graphics.plot_regress_exog(model4, 'price', fig=plt.figure(figsize=(12, 10)))
plt.show()

# Data transformation and visualization
data['age'] = abs(data['year'] - 2020)
data['engineSize_category'] = pd.Categorical(data['engineSize'])
sns.boxplot(data=data, x='engineSize_category', y='tax')
plt.xticks(rotation=90)
plt.show()
```

OLS Regression Results

```

=====
=====
Dep. Variable:          mpg    R-squared:
0.480
Model:                  OLS    Adj. R-squared:
0.480
Method:                 Least Squares    F-statistic:
9814.
Date:                   Sun, 06 Aug 2023    Prob (F-statistic):
0.00
Time:                   15:07:46    Log-Likelihood:           -3
7635.
No. Observations:      10617    AIC:              7.52
7e+04
Df Residuals:          10615    BIC:              7.52
9e+04
Df Model:               1
Covariance Type:        nonrobust
=====
=====

```

	coef	std err	t	P> t	[0.025	
0.975]						
const	67.2274	0.187	358.999	0.000	66.860	6
price	-0.0007	7.43e-06	-99.067	0.000	-0.001	-

```

=====
=====
Omnibus:                148.695    Durbin-Watson:
1.778
Prob(Omnibus):          0.000    Jarque-Bera (JB):           24
1.600
Skew:                   0.123    Prob(JB):                   3.4
5e-53
Kurtosis:               3.697    Cond. No.                   5.8
0e+04
=====
=====

```


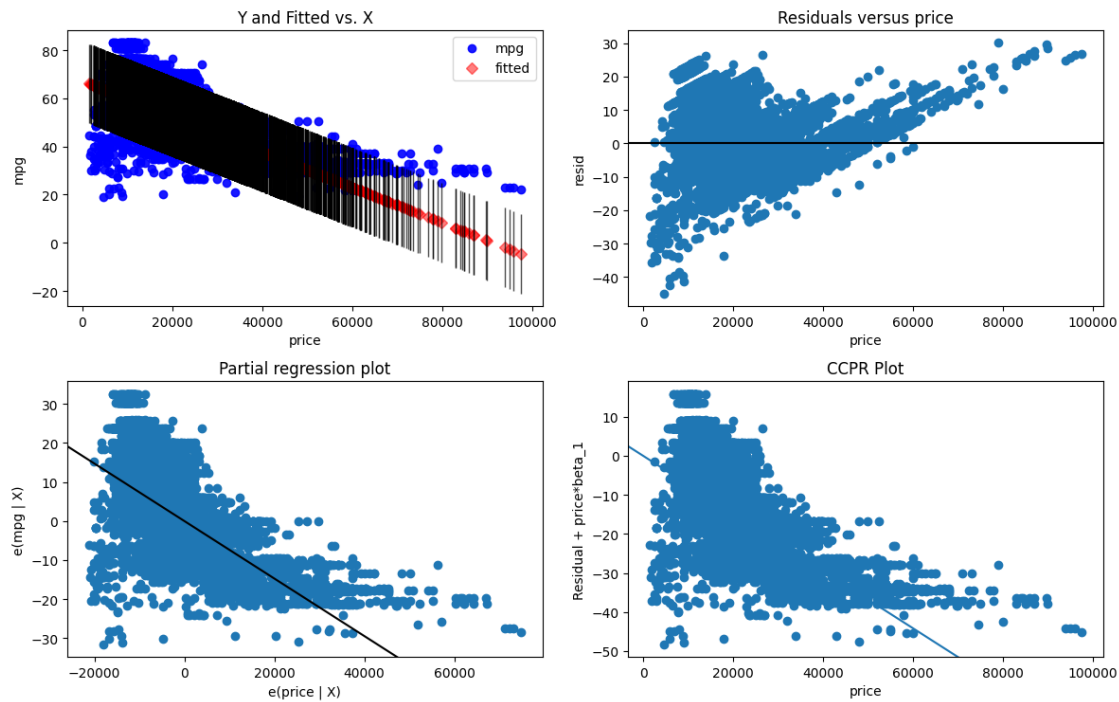
Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.8e+04. This might indicate that there are strong multicollinearity or other numerical problems.

eval_env: 1

Regression Plots for price



In []:

